

Muhammad Ghulam Jillani

[in LinkedIn](#) | [+92-321-1174167](#) | [+92-321-1179584](#) | [JillaniPortfolio](#) | [m.g.jillani123@gmail.com](#) | [Kaggle](#) | [GitHub](#) | [Medium](#)

Professional Summary

Lead AI Data Scientist and Machine Learning Engineer specializing in Generative AI, LLMs, and Autonomous AI Systems, with a proven record of transforming SaaS and PaaS platforms through innovative enterprise AI and agentic AI solutions. Expertise in optimizing workflows, streamlining data pipelines, and building scalable AI architectures to solve complex business challenges. Recognized as a 24x LinkedIn Top Voice, Top 100 Global Kaggle Master, and KaggleX BIPOC Mentor, contributing to the NVIDIA Developer Program, Google Developer Group, and AWS AI Community. Demonstrated leadership in AI-driven product innovation, MLOps and LLMOps strategies, and multimodal AI, delivering impactful results with cutting-edge technologies across industries.

Technical Skills

Programming & Development:

- Languages & Frameworks:** Python, Scikit-Learn, TensorFlow, Keras, PyTorch, NLTK, Hugging Face Transformers, OpenCV, FastAPI, Flask, Streamlit.
- Tools & Libraries:** Pandas, NumPy, Matplotlib, Plotly, Seaborn, PySpark, OpenAI API, REST APIs, GraphQL, Neo4j, Docker, GitHub Actions, CI/CD Pipelines.
- Database Systems:** SQL, NoSQL, Vector Databases (Pinecone, Faiss, Chroma DB, Weaviate, Qdrant).

Data Science & AI:

- Core Skills:** Machine Learning, Deep Learning, Generative AI, LLMs (GPT, Gemini, LLaMA, Falcon, DeepSeek), Natural Language Processing (NLP), Time Series Analysis, Model Deployment, Prompt Engineering.
- Frameworks & Tools:** LangChain, RAG (Retrieval-Augmented Generation), LlamaIndex, LangGraph, PhiData, LangServer, AutoGen, LangSmith, AutoML, N8N AI Agents.
- Applications:** AI-Driven Process Automation, Predictive Modeling, Statistical Analysis, Big Data Technologies, Data Visualization.

Cloud, MLOps & LLMOps:

- Platforms:** AWS (SageMaker, Bedrock, Lambda, EC2, Step Functions), Azure (Azure ML, Azure AI Studio, App Services, AKS), GCP (Vertex AI, Cloud Functions, BigQuery ML), Heroku, Vercel, and Hugging Face.
- Practices & Tools:** MLOps, LLMOps, AIOps, Cloud ML Pipelines, RAG, MLflow, LangChain, LangGraph, LlamaIndex, OpenLLM, CI/CD, Model Monitoring, and Orchestration Frameworks.

Management & Operations:

- Agile Methodologies, Microservices Architecture, Business Analysis, Product Management, Team Leadership, Stakeholder Management, Business Intelligence.

Communication & Collaboration:

- Strong in technical writing, Effective Communication, Stakeholder Engagement, and Team Collaboration, Fostering a Productive and Inclusive Environment.

Languages:

- Fluent in English.

Experience

Senior Data Scientist & Machine Learning Software Engineer (Generative AI)

PURELOGICS

New York, United States12/2024 – Currently

- AI-Driven MVPs & Enterprise Deployments:** Led the development of high-impact AI solutions, including an Enterprise-Grade RAG Pipeline for Business Intelligence, a Health Guard AI System for Remote Patient Monitoring, a GPT-4o-Based Conversational AI Chatbot, and an Agentic AI System for Legal & Financial Advisory. Designed and deployed AI-Powered Meeting Assistant (AIMA) to streamline meeting management using Generative AI, RAG, and Autonomous AI Agents.
- Scalable AI & Cloud Solutions:** Architected and deployed end-to-end AI/ML pipelines on AWS (SageMaker, Bedrock, Lambda, API Gateway) and GCP (Vertex AI, BigQuery, Cloud Run). Ensured high availability, cost efficiency, and real-time insights for AI-driven applications.
- Advanced Multi-Agent AI Architectures:** Engineered multi-agent LLM systems with LangChain, LangGraph, Pinecone, Graph RAG, and Knowledge Graphs to enable autonomous decision-making, automation, and AI-driven workflows across industries.
- Leadership & AI Strategy Execution:** Led cross-functional teams, aligning AI strategies with business objectives and optimizing LLMOps, MLOps, and AIOps pipelines for seamless model lifecycle management, from development to enterprise deployment.
- Mentorship & AI Innovation:** Mentored AI/ML teams, fostering expertise in LLM fine-tuning, Retrieval-Augmented Generation (RAG), and AI-powered automation. Drove continuous innovation in Generative AI applications, enhancing AI-driven productivity and operational efficiency.
- Key Achievements:** Successfully delivered scalable, enterprise-grade AI solutions, exceeding performance benchmarks, accelerating deployment cycles, and enhancing business intelligence, automation, and AI-driven decision-making.

Senior Data Scientist Machine Learning & Gen AI Engineer (Volunteer, Part-Time – 20 hrs/week)

NVIDIA

California, United States08/2024 – Currently

- NVIDIA & Z by HP Developer Programs:** Active member contributing to GPU-optimized applications and testing enterprise-level Generative AI solutions on NVIDIA AI Studio and Z by HP platforms.
- LLM Contributions:** Developed and tested open-source and enterprise-level LLMs, including vLLMs, for Generative AI domains.
- Specialized SDK Integration:** Leveraged cutting-edge NVIDIA tools and SDKs to build and optimize AI solutions.
- Collaborative Expertise:** Worked closely with NVIDIA and Z by HP teams to validate and enhance AI tools and frameworks.
- Community Contributions:** Actively contributed to open-source LLM projects, driving innovation in Generative AI and large-scale deployments.

Senior Data Scientist & Machine Learning Engineer

BLOCBELT

Maryland, United States12/2022 – 12/2024

- Spearheaded AI-driven SaaS and PaaS development, leading a team of three in Big Data, AI Engineering, and Data Analysis—boosting operational efficiency and saving \$2M for clients.
- Improved sales forecasting accuracy by 80–90% across 120+ market entities, driving data-backed strategies and enhanced workflows.
- Launched a personalized fashion e-commerce platform with a recommendation engine, increasing sales by 80% and online orders by 200% ahead of schedule.
- Streamlined data workflows and **built image processing pipelines using OpenCV**, significantly reducing processing costs and boosting automation efficiency.
- Applied **computer vision techniques for product image analysis**, enhancing catalog accuracy and digital asset classification.
- Delivered an 18% revenue increase through predictive analytics and statistical modeling, with a 45% improvement in data visibility via custom dashboards.
- Executed over 25 A/B tests to optimize digital experiences and **deployed image-driven insights into the AI-powered shopping journey**.

- **Championed** mentorship initiatives through the Kaggle-X BIPOC Mentorship Program, dedicating 20 hours weekly to guiding aspiring data scientists in technical skill development and community contributions.
- **Secured** a position among the Top 100 Kaggle contributors globally, recognized as the first Pakistani mentor in the program.
- **Led** a team of four mentees, enhancing their skills in data science, machine learning, and real-world problem-solving.
- **Drove** innovation by collaborating on high-impact AI projects and actively contributing to the global data science community.

Data Scientist
Thailand

Crypto-Express
01/2022 - 12/2022

- Developed an Anti-spoofing Face-App, enhancing digital identity security and reducing identity fraud by 70%.
- Applied AI and ML techniques to solve industry-critical issues, improving efficiency by 25% and reducing operational costs by 15%.
- Engineered predictive models and provided actionable insights, leading to a 20% revenue increase through data-driven decision-making.
- Advanced expertise in EDA, Machine Learning, and Deep Learning, saving over 250+ hours of manual analysis annually.
- Successfully deployed AI/ML projects remotely, adapting to dynamic conditions and gaining significant experience in the AI/ML field.

Artificial Intelligence Engineer
Lahore, Pakistan

Pakistan Freelancing Training Center
01/2021 - 12/2021

- Led AI and Data Science training sessions, preparing students to tackle real-world challenges.
- Designed and deployed models using TensorFlow and Keras, improving system efficiency and accuracy.
- Bridged the gap between academic learning and industry skills, driving practical applications of AI technologies.
- Automated repetitive data tasks using Python, significantly reducing manual labor and error rates.
- Engaged in ongoing AI research to apply cutting-edge methodologies to projects.

Projects

- **Enterprise Knowledge Navigator – Agentic RAG for Corporate Wikis:** Built an AI-powered internal assistant for navigating vast enterprise documentation across departments. Combined multi-agent reasoning, metadata tagging, and vector-enhanced retrieval to help employees access policies, SOPs, HR guidelines, and technical manuals.
- **LegalDoc AI – Intelligent Document Retrieval & Summarization System:** A secure enterprise platform that uses Agentic RAG to analyze legal contracts and case files. Integrated an LLM-powered chatbot for dynamic querying and document-based reasoning across multi-format legal datasets.
- **Healthcare AI Assistant (LLMOps, Bedrock):** Developed an advanced LLMOps pipeline for a healthcare AI assistant using Amazon Bedrock, fine-tuned Gemini 1.5, and SageMaker Pipelines. Integrated automated dataset versioning and continuous monitoring using Evidently AI. Achieved 85% accuracy in medical diagnosis recommendations, streamlining workflows for healthcare professionals.
- **AI-Powered Blog Generator (LLM):** Developed an AI-powered content generator leveraging Llama 3.1 8B, AWS Bedrock, and RAG for real-time fine-tuning on dynamic datasets. Integrated with Streamlit for a seamless user interface, the system automated blog creation, editing, and publishing, increasing efficiency in content marketing and engagement for enterprises.
- **Conversational AI Chatbot for Customer Support (GPT-4o, LangChain, Bedrock):** Designed and deployed an enterprise-grade chatbot powered by GPT-4o, integrated with LangChain and AWS Bedrock. Implemented advanced NLP techniques such as intent recognition and dynamic response generation, reducing customer query resolution times by 40% and improving customer satisfaction.
- **AutoML-Studio (MLOps):** Built a low-code/no-code machine learning platform enabling rapid development and deployment of predictive models. Incorporated MLflow for tracking, XAI (SHAP) for explainability, and data drift detection to ensure model reliability. Deployed with FastAPI and Streamlit, the platform supports classification, regression, time-series forecasting, and clustering, empowering non-technical users to build robust AI solutions.
- **AI-Powered Meeting Assistant (LLMOps):** Developed an AI-driven Meeting Management System using Streamlit, FastAPI, LangChain, OpenAI GPT-4o, and Pinecone for intelligent meeting transcription, summarization, and retrieval. Integrated RAG-based retrieval, Google Sheets API for roadmap estimations, and automated email dispatch via Gmail API. Deployed the system on AWS for scalability, enabling real-time conversation analysis and enterprise-level knowledge retrieval.
- **DocuWiz AI (LLMOps):** Engineered a document intelligence tool leveraging advanced NLP models like BART and DistilBERT for text extraction, summarization, and sentiment analysis. Integrated RAG and LLMs to enable detailed insights and visualizations for PDF and DOCX analysis. Deployed as an interactive Streamlit application and FastAPI, reducing document processing time by 60% and enhancing productivity in legal and academic domains.
- **Harvestify-AI-Powered-Plant-Health-Assistant (AIOps):** Designed and implemented an AI-driven application leveraging Azure AI Studio and Azure App Services to diagnose plant leaf diseases through real-time image analysis. The system integrates deep learning models to recommend crops and fertilizer solutions tailored to soil and weather conditions, enhancing agricultural decision-making. Deployed on Microsoft Azure, the solution reduced disease detection time by 70%, improved crop yield predictions by 30%, and empowered farmers with actionable insights to optimize productivity and sustainability.
- **Intelligent Document Summarization Tool (LLMOps):** Developed an Intelligent Document Summarization Tool using transformer-based LLM models like Pegasus to generate concise summaries from lengthy documents while preserving key information. Deployed with FastAPI and Streamlit APP for real-time interaction, supporting formats like PDF and DOCX, and leveraging semantic analysis for relevance extraction. Reduced manual document review time by 60%, boosting productivity in legal and academic workflows.
- **Customer Satisfaction Prediction System (MLOps):** Created a predictive analytics system using ZenML and MLflow within an MLOps framework to achieve 93% accuracy in forecasting customer satisfaction. Streamlined pipeline management and deployed models efficiently, providing actionable insights that improved business decision-making and customer retention strategies.
- **AIOps-Driven Predictive Maintenance System:** Created an AIOps-based predictive maintenance platform leveraging AWS Sage Maker, Kafka, and Databricks. Implemented real-time anomaly detection using LSTMs and automated model retraining to prevent equipment failures. Deployed in a cloud-native environment with CI/CD pipelines for continuous updates, improving maintenance efficiency by 40% and reducing downtime.

Education

Bachelor in Computer Science (Major: Artificial Intelligence and Computer Science)

Institute of Management Sciences

CERTIFICATIONS

- **IBM Machine Learning Specialization Professional Certificate**, IBM.
 - **Microsoft Azure AI Fundamentals AI-900 Exam Prep Specialization**, Microsoft.
 - **Machine Learning Engineering for Production (MLOps)**, DeepLearning.ai.
 - **Generative Adversarial Networks (GANs) Specialization**, DeepLearning.ai.
 - **Preparing for Google Cloud Certification: Machine Learning Engineer**, Google.
 - **AWS Cloud Solutions Architect Professional Certificate**, Amazon Web Services.
 - **Practical Data Science on the AWS Cloud Specialization**, DeepLearning.ai.
 - **MLOps | Machine Learning Operations Specialization**, Duke University.
 - **Advanced Machine Learning on Google Cloud Specialization**, Google Cloud.
 - **Google Advanced Data Analytics Professional Certificate**, Google.
- **Deep Learning Specialization**, DeepLearning.ai.
 - **Prompt Engineering for ChatGPT**, Vanderbilt University.
 - **AI Product Management**, Duke University.
 - **IBM Generative AI Product Managers**, IBM.
 - **IBM AI Product Manager**, IBM.
 - **Google Project Management: Professional Certificate**, Google.
 - **Machine Learning Specialization**, DeepLearning.AI.
 - **Google Business Intelligence Professional Certificate**, Google.
 - **Large Language Model Operations (LLMOps)**, Duke University.